

# Spatial indexes for OSM in PostGIS

<https://slides.com/fxku/sotm19>

Felix Kunde



```
CREATE INDEX pts_spx  
ON point_table  
USING GIST (geom)
```

```
CREATE INDEX pts_spx  
ON point_table  
USING GIST (geom)  
BRIN < v2.3
```


```
CREATE INDEX pts_spx  
ON point_table  
USING GIST (geom)
```

BRIN < v2.3

SPGIST < v2.5


# GiST

It's a framework  
PostGIS > R-Tree  
2-nD and kNN support




Leaves





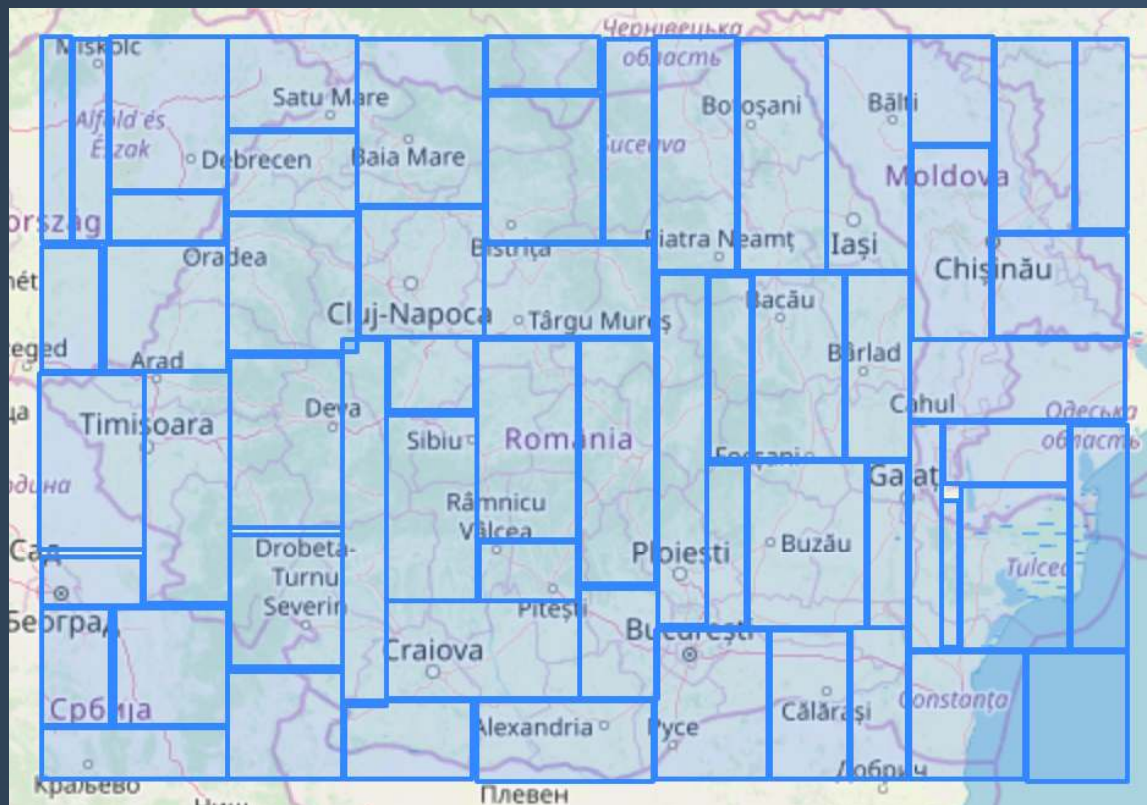


Nodes









Level 1



Level 2

Produced with Gevel

# ST\_Intersects in ms

tests	100 k	1 Mio	10 Mio	100 Mio	1 Bn
no index	18.00	87.00	670	6473	135529
bulk	0.21	14.00	19	146	1568
online	0.18	15.00	29	163	1672
vacuum	0.16	0.87	18	145	1551
cluster	0.13	0.64	16	32	214

# ST\_Intersects in ms

tests	100 k	1 Mio	10 Mio	100 Mio	1 Bn
no index	18.00	87.00	670	6473	135529
bulk	0.21	14.00	19	146	1568
online	0.18	15.00	29	163	1672
vacuum	0.16	0.87	18	145	1551
cluster	0.13	0.64	16	32	214

# ST\_Intersects in ms

tests	100 k	1 Mio	10 Mio	100 Mio	1 Bn
no index	18.00	87.00	670	6473	135529
bulk	0.21	14.00	19	146	1568
online	0.18	15.00	29	163	1672
vacuum	0.16	0.87	18	145	1551
cluster	0.13	0.64	16	32	214

# ST\_Intersects in ms

tests	100 k	1 Mio	10 Mio	100 Mio	1 Bn
no index	18.00	87.00	670	6473	135529
bulk	0.21	14.00	19	146	1568
online	0.18	15.00	29	163	1672
vacuum	0.16	0.87	18	145	1551
cluster	0.13	0.64	16	32	214

# BRIN



Block ranges. That's all.  
Data should be sorted!  
PostGIS > BBox per block  
2-4D, no kNN



# Block Range Index examples

...	time	geom
...	2019-03-13	POINT(13.8 50)
...	2019-03-13	POINT(13.8 50)
	...	...
	2019-03-14	POINT(13.8 51)
	2019-03-16	POINT(13.7 51)
	...	...
...	2019-03-16	POINT(13.7 51)
...	2019-03-16	POINT(13.7 51)

**(2019-03-13 09:00:00 , 2019-03-14 11:00:00)**

**(A-Weg , Grunaer Str.)**

**BBOX**

# Block Range Index examples

...	time	geom
...	2019-03-13	POINT(13.8 50)
...	2019-03-13	POINT(13.8 50)
	...	...
	2019-03-14	POINT(13.8 51)
	2019-03-16	POINT(13.7 51)
	...	...
...	2019-03-16	POINT(13.7 51)
...	2019-03-16	POINT(13.7 51)

**(2019-03-13 09:00:00 , 2019-03-14 11:00:00)**

**(A-Weg , Grunaer Str.)**

**BBOX**

**(2019-03-14 11:15:00 , 2019-03-15 18:00:00)**

**(Grunaer Weg, Nürnberger Str.)**

**BBOX**

# Block Range Index examples

...	time	geom
...	2019-03-13	POINT(13.8 50)
...	2019-03-13	POINT(13.8 50)
	...	...
	2019-03-14	POINT(13.8 51)
	2019-03-16	POINT(13.7 51)
	...	...
...	2019-03-16	POINT(13.7 51)
...	2019-03-16	POINT(13.7 51)

**(2019-03-13 09:00:00 , 2019-03-14 11:00:00)**

**(A-Weg , Grunaer Str.)**

**BBOX**

**(2019-03-14 11:15:00 , 2019-03-15 18:00:00)**


**(Grunaer Weg, Nürnberger Str.)**

**BBOX**

**(2019-03-16 09:00:00 , 2019-03-16 17:30:00)**

**(Oberauer Str. , Zwinglistr.)**

**BBOX**







**ORDER BY geom**  
(PostGIS v3.0 uses Hilbert Curve)



**ORDER BY**  
**ST\_GeoHash(geom)**

```
SET enable_seqscan = false;
```

*(discourage Postgres' query planner to use seq scans)*

# BRIN vs. GiST


tests	100 k	1 Mio	10 Mio	100 Mio	1 Bn
create gist	700 ms	8 sec	2 min	23 min	6 hrs
create brin	24 ms	0.2 sec	2 sec	18 sec	90 sec
size gist	5 MB	50 MB	500 MB	5 GB	50 GB
size brin	24 KB	24 KB	48 KB	376 KB	3,6 MB
duration	x25	x23	x1,4	x1.6	x1.1

# sp-GiST



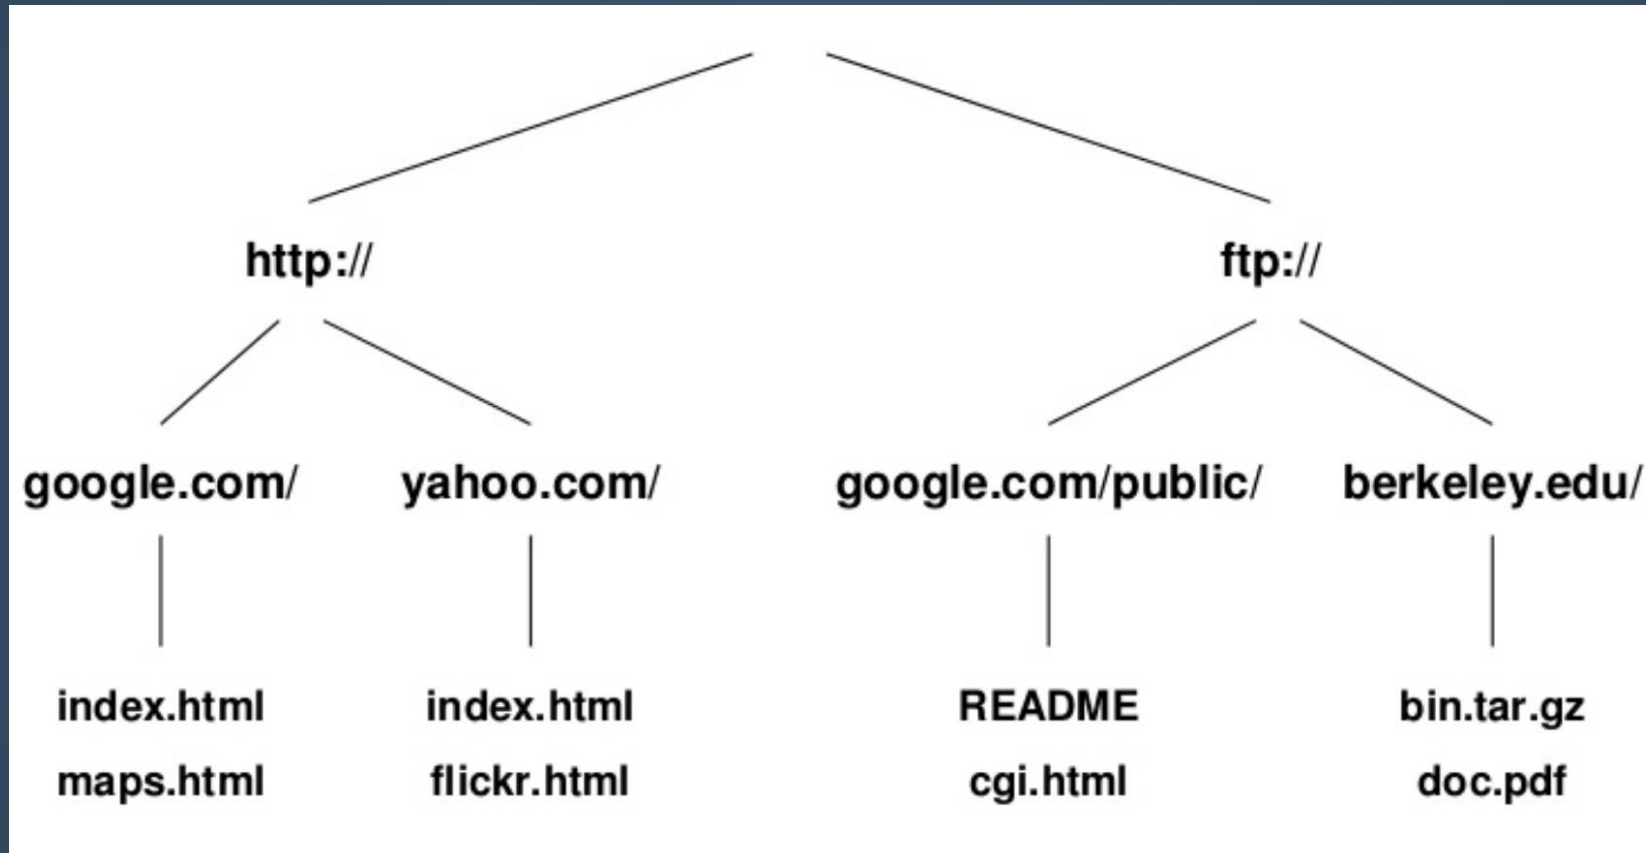
Framework like GIST  
Unbalanced tree  
No overlaps & prefixes  
2-nD, no kNN, yet

# kd-Tree, Quadtree




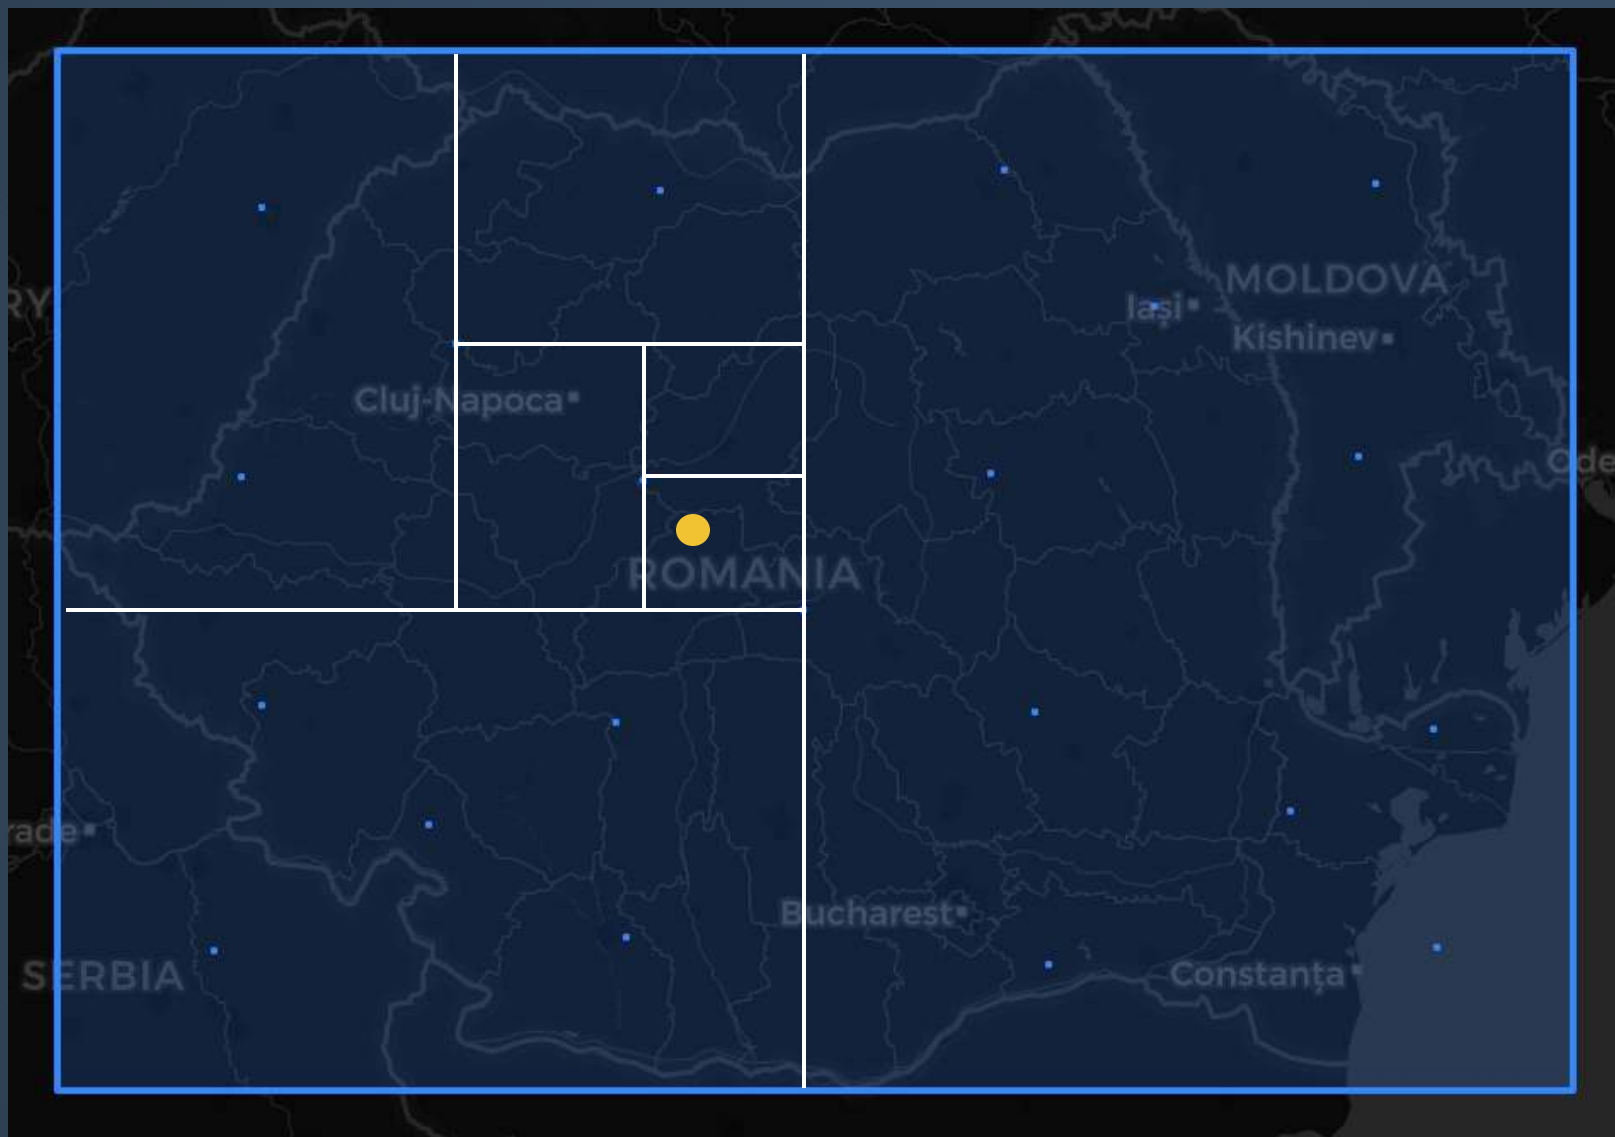
[https://www.researchgate.net/figure/Adaptive-k-d-tree\\_fig9\\_2334587](https://www.researchgate.net/figure/Adaptive-k-d-tree_fig9_2334587)

# Split values between index and leaf nodes



Momjian 2019: Flexible Indexing with Postgres [\[PDF\]](#)

# Trick: No Overlap via more dimensions



Each point you see on the map are in fact 4 bounding boxes which are the prefixes of the sp-GiST tree defining the bounds of child quadrants

# sp-GiST vs. GiST

tests	100 k	1 Mio	10 Mio	100 Mio	1 Bn
create gist	700 ms	8 sec	2 min	23 min	6 hrs
create spgist	344 ms	3,7 sec	50 sec	11 min	8 hrs
size gist	5 MB	50 MB	500 MB	5 GB	50 GB
size spgist	4,5 MB	44 MB	440 MB	4.3 GB	43 GB
duration	x0.85	x1	x1	x1	x1.1

# Thematic filters

# Multi-column index

```
CREATE INDEX idx ON planet_osm_point  
USING GIST (way, power);
```

Requires btree\_gist extension

Spatial column first as it is more selective

Less read from heap

Index will be bigger, combine with ...

# Partial index

```
CREATE INDEX idx ON planet_osm_point  
USING GIST (way)  
WHERE power = 'pole';
```

Smaller index, less IO overhead

Very specific to a certain task

OSM people **do** this



# Covering index

```
CREATE INDEX idx ON planet_osm_point  
USING GIST (way)  
INCLUDE (power);
```

"Payload" only in the leaves (not sorted)  
No big speed difference to multi-column index  
GiST support in Postgres 12

## Btw: Index-only scans?

Not possible due to **compression**

Requires new index tuple format with SRID

For small geoms (Points, BBox, ..?)

Breaking all GiSTs out there?

# GeoHash index

```
CREATE INDEX idx ON planet_osm_point
USING BTREE (
    ST_GeoHash(ST_Transform(way, 4326)), power
)
WHERE power IS NOT NULL;
```

BTREE benefits (parallel build, index-only-scan)

Spatial filtering tricky (GeoHash grid)

Slower, so only useful for table clustering

# Radix tree

```
CREATE INDEX idx ON planet_osm_point
USING SPGIST (
    ST_GeoHash(ST_Transform(way, 4326)) text_ops
)
WHERE power IS NOT NULL;
```






GeoHash and SPGIST seem like a great fit  
LIKE not supported, but can use ranges  
Slow (only two prefixes on small dataset)

# Postgres geometry


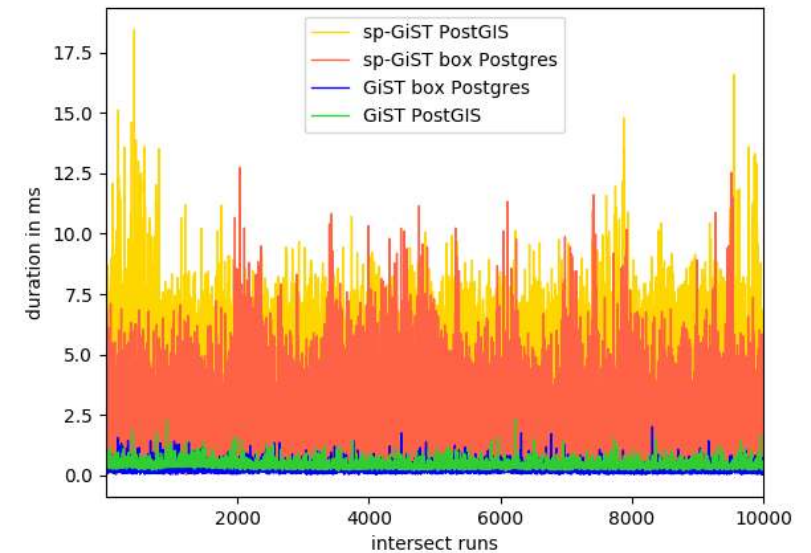

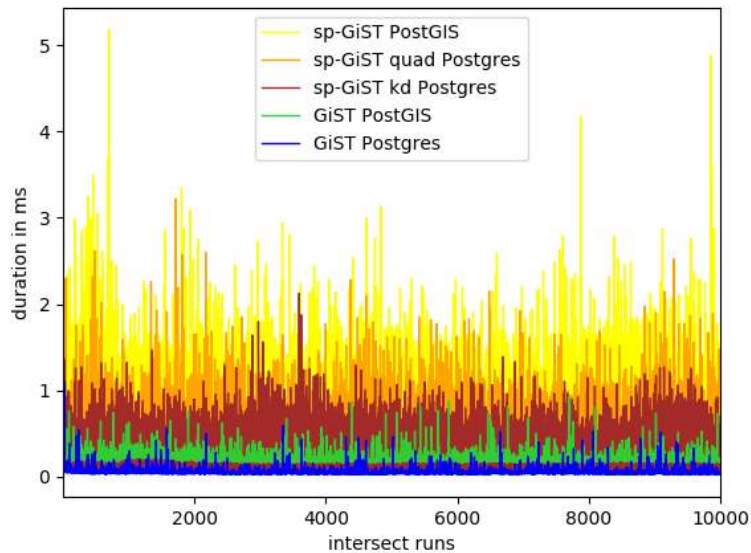
```
CREATE INDEX idx ON planet_osm_point
USING SPGIST (
    CAST (way AS point) quad_point_ops
)
WHERE power IS NOT NULL;
```

Queries need casting but it works  
And performance is good (smaller tuple layout)

```
SELECT 1 FROM planet_osm_[point|line|polygon]
WHERE CAST(way AS point) <@ [filter-box]
AND _ST_Intersects(way, [filter-buffer]);
```






```
SELECT 1 FROM planet_osm_[point|line|polygon]
WHERE CAST(way AS point) <@ [filter-box];
```




# OSM Experiments



			
Points	0.5 Mio, 128 MB, <b>12%</b> 132, 53, 0.2 MB	1.7 Mio, 0.4 GB, <b>13%</b> 96, 92, 0.5 MB	11.7 Mio, 3 GB, <b>11%</b> 630, 600, 0.1 MB
Lines	0.5 Mio, 250 MB, <b>43%</b> 132, 53, 0.2 MB	2 Mio, 1 GB, <b>44%</b> 230, 93, 0.6 MB	14.37 Mio, 7 GB, <b>42%</b> 1500, 600, 0.2 MB
Polygons	1.2 Mio, 533 MB, <b>38%</b> 132, 53, 0.2 MB	5 Mio, 2.3 GB, <b>39%</b> 560, 220, 0.2 MB	35.6 Mio, 16 GB, <b>39%</b> 3800, 1500, 0,5 MB

*Rowcount, Tablesize, Geom. size to table size  
sizes of GiST, sp-GiST and BRIN*



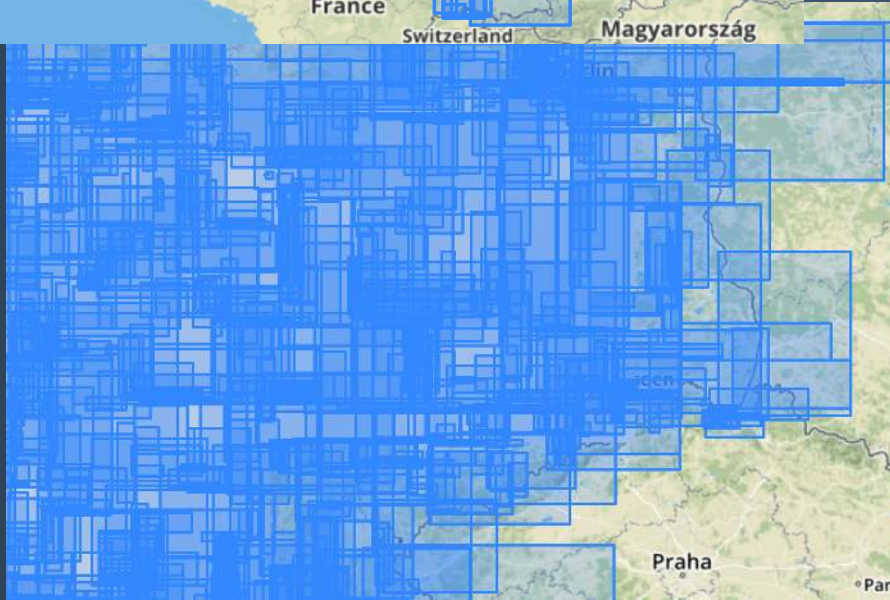
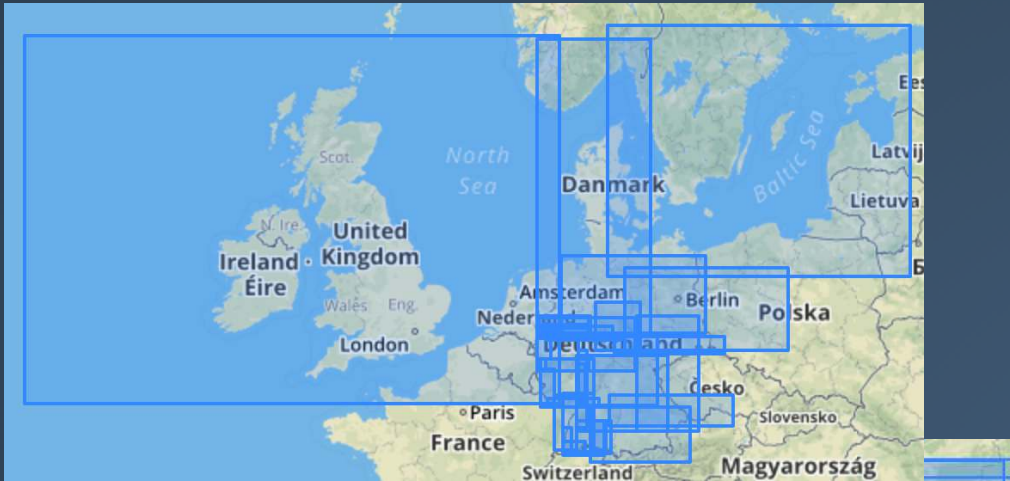
Kudos to **DBeaver**



GiST Index on  
`planet_osm_point`







# GiST Index on planet\_osm\_line







**ST\_Subdivide me!**

# GiST Index on planet\_osm\_polygon









The more data, the closer the gap to BRIN  
GiST and sp-GiST comparable on Polygons



highway:  
emergency\_access\_point  
~ 25.000



highway: motorway  
~ 65.000



natural: bare\_rock  
~ 25.000



sp-GiST slightly faster for points and polygons  
Higher hit-rate of shared buffers (when is this bad?)

```
CREATE INDEX highway_eap_gist ON planet_osm_point USING GIST (way) WHERE highway = 'emergency_access_point';
CREATE INDEX highway_eap_spgist ON planet_osm_point USING SPGIST (way) WHERE highway = 'emergency_access_point';
```

```
EXPLAIN (ANALYZE, Buffers)
SELECT 1 FROM planet_osm_point, random_points
WHERE ST_Intersects(way, ST_Buffer(geom, 1000))
AND highway = 'emergency_access_point';
```

#### GIST

```
Nested Loop (cost=0.28..29261.90 rows=83000 width=4) (actual time=0.685..452.657 rows=817 loops=1)
  Buffers: shared hit=21756
  -> Seq Scan on random_points (cost=0.00..94.00 rows=10000 width=32) (actual time=0.011..1.977 rows=10000 loops=1)
    Buffers: shared hit=84
  -> Index Scan using highway_eap_gist on planet_osm_point (cost=0.28..2.92 rows=1 width=32) (actual time=0.018..0.019 rows=0 loops=10000)
    Index Cond: (way && st_buffer(random_points.geom, '1000'::double precision))
    Filter: _st_intersects(way, st_buffer(random_points.geom, '1000'::double precision))
    Rows Removed by Filter: 0
    Buffers: shared hit=21672
Planning Time: 0.201 ms
Execution Time: 452.791 ms
```

#### sp-GIST

```
Nested Loop (cost=0.28..29274.00 rows=83000 width=4) (actual time=1.267..362.863 rows=817 loops=1)
  Buffers: shared hit=39917
  -> Seq Scan on random_points (cost=0.00..94.00 rows=10000 width=32) (actual time=0.019..1.690 rows=10000 loops=1)
    Buffers: shared hit=84
  -> Index Scan using highway_eap_spgist on planet_osm_point (cost=0.28..2.92 rows=1 width=32) (actual time=0.011..0.012 rows=0 loops=10000)
    Index Cond: (way && st_buffer(random_points.geom, '1000'::double precision))
    Filter: _st_intersects(way, st_buffer(random_points.geom, '1000'::double precision))
    Rows Removed by Filter: 0
    Buffers: shared hit=39833
Planning Time: 0.347 ms
Execution Time: 363.027 ms
```

```
CREATE INDEX motorway_gist ON planet_osm_line USING GIST (way) WHERE highway = 'motorway';
CREATE INDEX motorway_spgist ON planet_osm_line USING SPGIST (way) WHERE highway = 'motorway';
```

```
EXPLAIN (ANALYZE, Buffers)
SELECT 1 FROM planet_osm_line, random_points
WHERE ST_Intersects(way, ST_Buffer(geom, 1000))
AND highway = 'motorway';
```

#### GIST

```
Nested Loop (cost=0.28..93512.50 rows=229587 width=4) (actual time=3.490..457.310 rows=3023 loops=1)
  Buffers: shared hit=28119
  -> Seq Scan on random_points (cost=0.00..94.00 rows=10000 width=32) (actual time=0.007..1.564 rows=10000 loops=1)
    Buffers: shared hit=84
  -> Index Scan using motorway_gist on planet_osm_line (cost=0.28..9.34 rows=2 width=218) (actual time=0.016..0.024 rows=0 loops=10000)
    Index Cond: (way && st_buffer(random_points.geom, '1000'::double precision))
    Filter: _st_intersects(way, st_buffer(random_points.geom, '1000'::double precision))
    Rows Removed by Filter: 0
    Buffers: shared hit=28035
Planning Time: 0.278 ms
Execution Time: 457.651 ms
```

#### sp-GIST

```
Nested Loop (cost=0.28..93497.10 rows=229587 width=4) (actual time=8.248..653.375 rows=3023 loops=1)
  Buffers: shared hit=294815
  -> Seq Scan on random_points (cost=0.00..94.00 rows=10000 width=32) (actual time=0.010..1.836 rows=10000 loops=1)
    Buffers: shared hit=84
  -> Index Scan using motorway_spgist on planet_osm_line (cost=0.28..9.34 rows=2 width=218) (actual time=0.032..0.041 rows=0 loops=10000)
    Index Cond: (way && st_buffer(random_points.geom, '1000'::double precision))
    Filter: _st_intersects(way, st_buffer(random_points.geom, '1000'::double precision))
    Rows Removed by Filter: 0
    Buffers: shared hit=294731
Planning Time: 0.345 ms
Execution Time: 653.781 ms
```



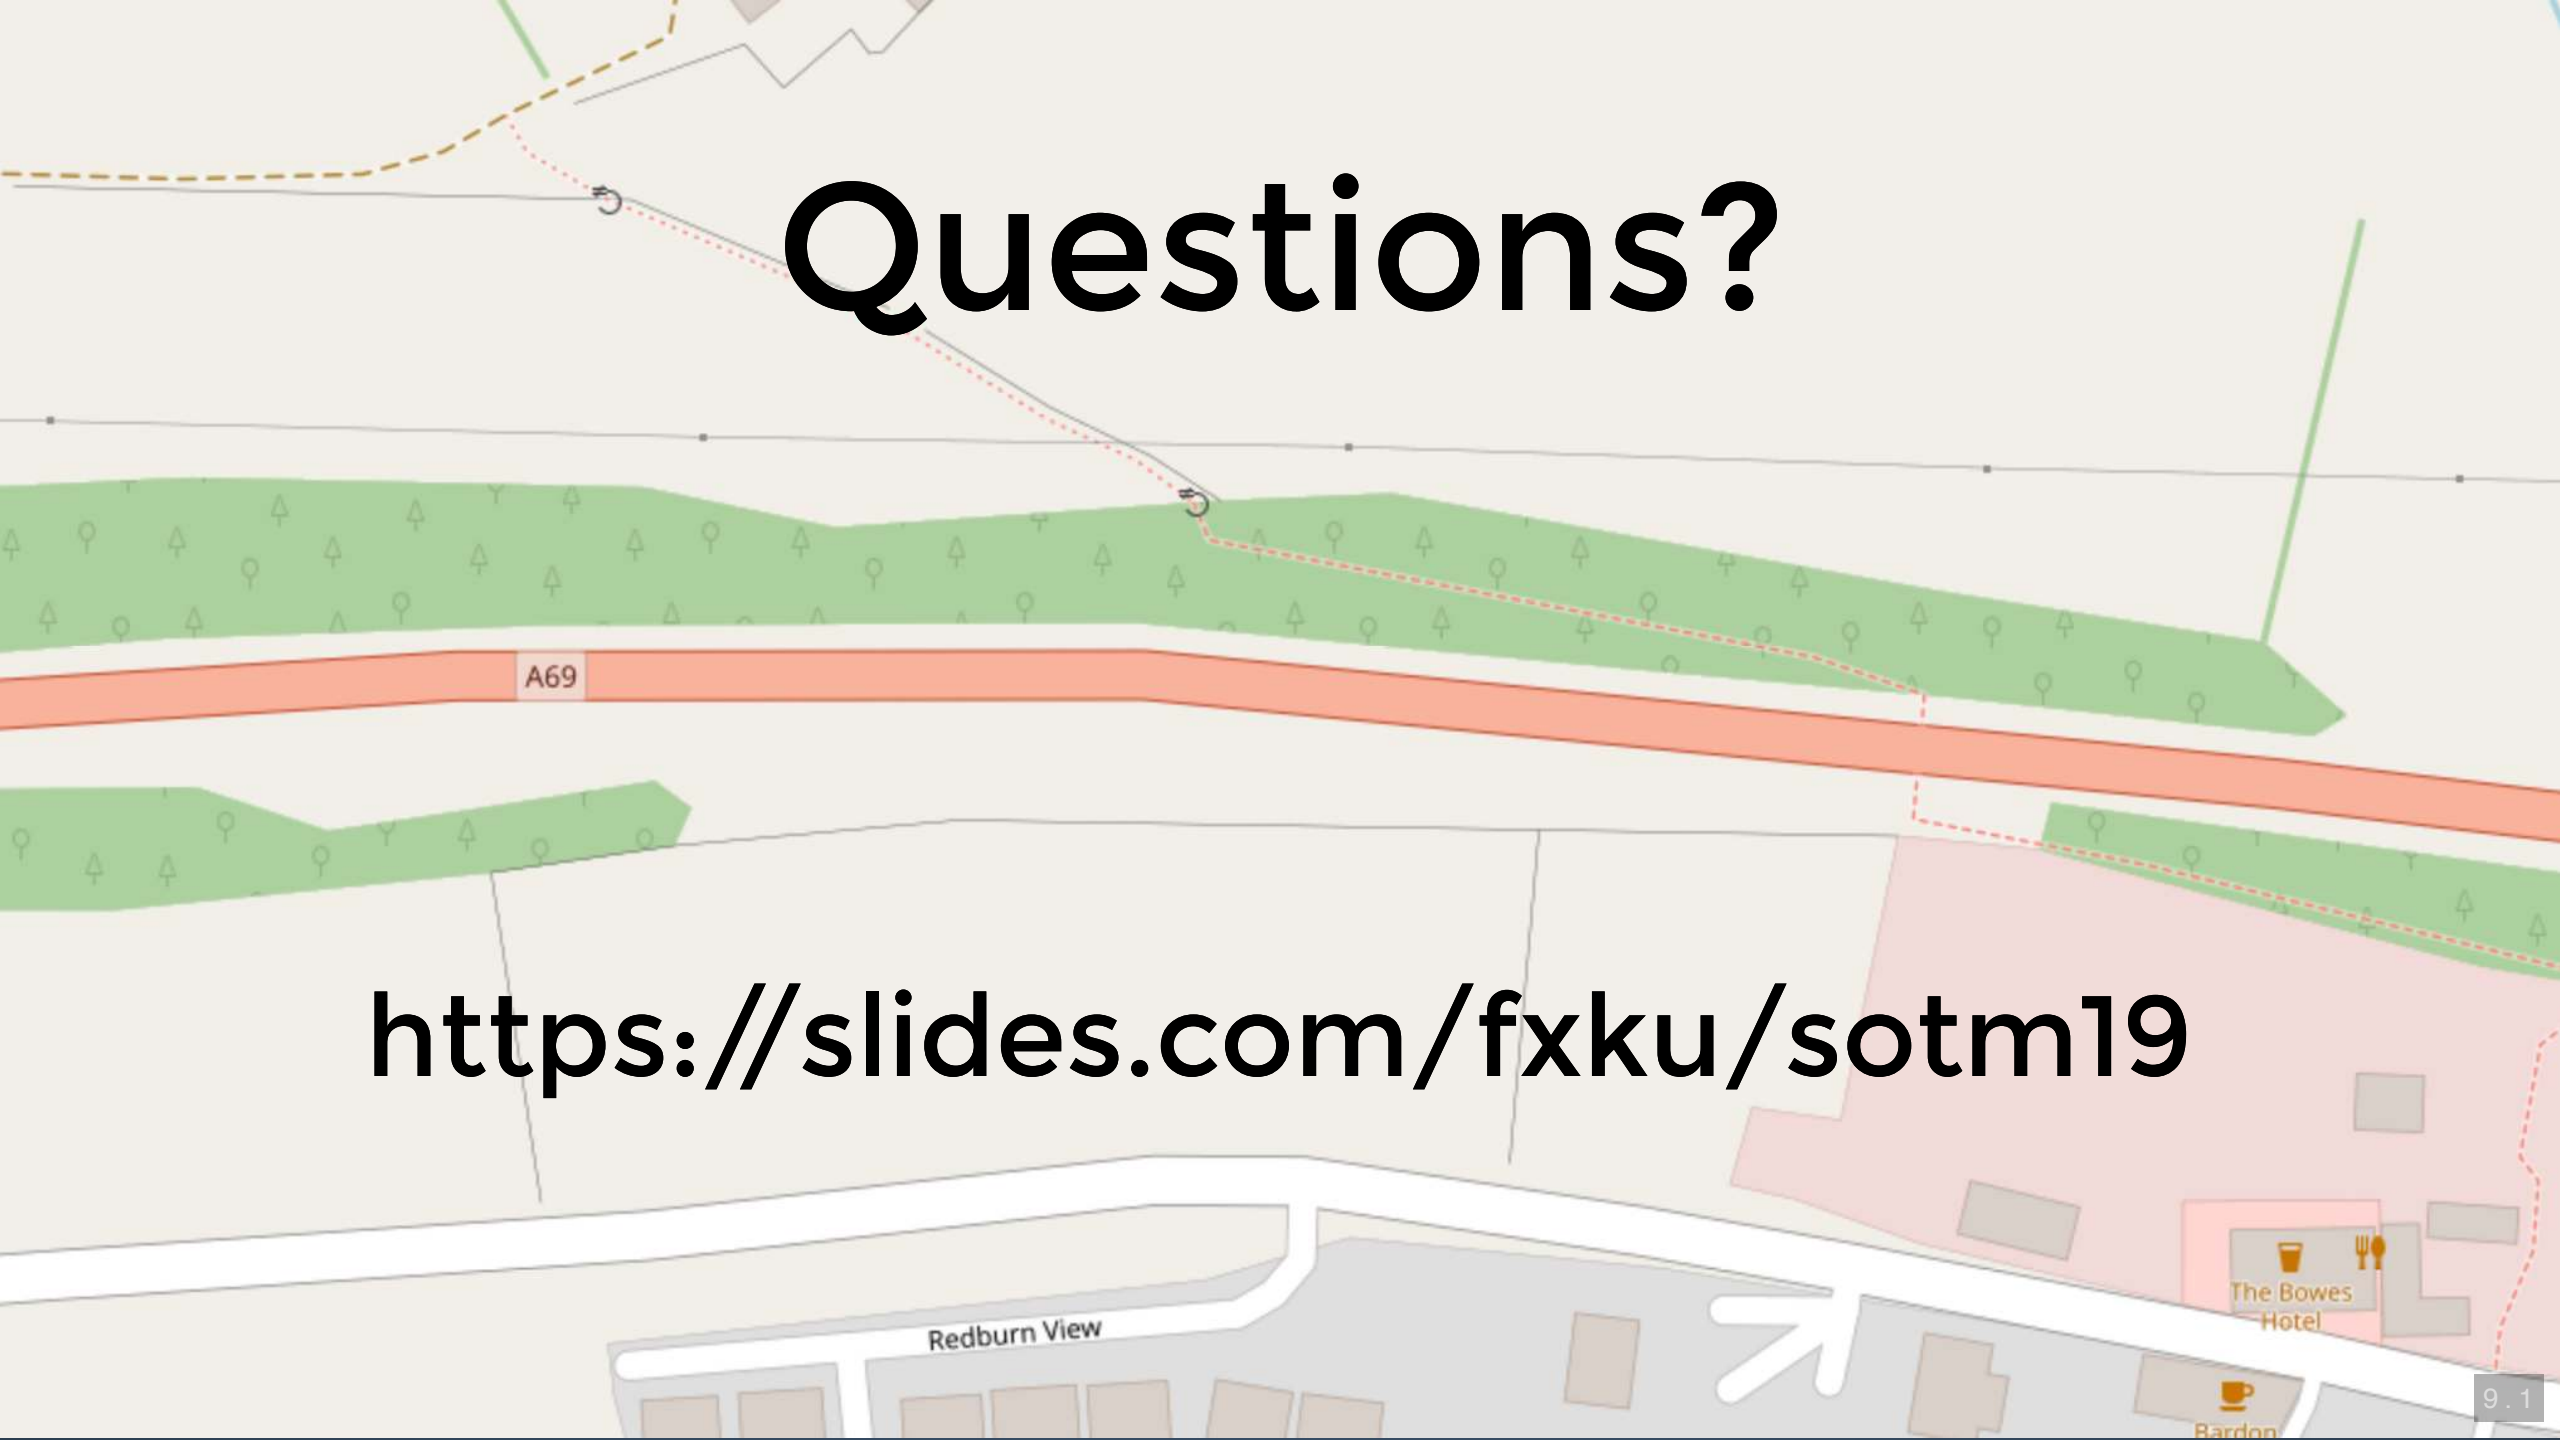
# Conclusion & Outlook

Stick with GiST  
But give BRIN a try  
Use partial indexes  
sp-GiST for unbalanced themes  
Keep your table stats up-to-date

**Parallel** queries!  
**Faster** index building!  
Index-only scans?  
**New** table access methods?

# Questions?

<https://slides.com/fxku/sotm19>



## Used hardware

Tuxedo Infinity Book 13

Intel i7-8550U CPU 1.80GHz

Quadcore, 8 CPUs

32 GB RAM

500GB SSD disk

## Repo

[github.com/FxKu/postgis\\_indexing](https://github.com/FxKu/postgis_indexing)

## PostgreSQL config

PostgreSQL 11 & PostGIS 2.5

shared\_buffers = 16 GB

work\_mem = 128 MB

maintenance\_work\_mem = 4 GB

min/max\_wal\_level = 16/4 GB

checkpoint\_timeout = 30 min

checkpoint\_completion\_target = 0.9

random\_page\_cost = 1.1

cpu\_tuple\_cost = 0.001

cpu\_index\_tuple\_cost = 0.001

effective\_cache\_size = 24 GB

default\_statistics\_target = 500

# More links

Pettus 2019: Look it up - Practical PostgreSQL indexing [\[PDF\]](#)

Rogov 2019: Indexes in PostgreSQL - sp-GiST ([Blog](#))

Katz 2018: Why covering indexes are helpful ([Blog](#))

Albe 2019: Optimizer support for functions ([Blog](#))

Wienand 2019: Be inclusive (Covering indexes) [\[PDF\]](#)