


# OSM Data Processing with PostgreSQL / PostGIS


Jochen Topf  
[jochentopf.com](http://jochentopf.com)



**OpenStreetMap**



**PostgreSQL**



osm2pgsql builds ways table but not line table

osm2pgsql

19 Jun, 08:48 SimonPoole ♦ 38.2k

how to get way OSM ID lat, lon from postgresql database?

psql

imposm

postgresql

postgres

21 Jan, 07:01 SimonPoole ♦ 38.2k

Update osm2pgsql database from apidb database

apidb

osm2pgsql

osmosis

14 Aug, 11:02 YUNtee 16

Nominatim PgSQL database size growing gradually

database

nominatim

size

29 Jul '16, 12:36 jot 496

osm2pgsql processing a 68MB planet diff for more than 19 hours

planet


diff

nominatim

osm2pgsql


updates

06 Jun, 10:36 SimonPoole ♦ 38.2k



# What we will talk about...

- Background: Relational Databases, Geodata
- Converting OSM Data
- Use Cases
- Tools
- Tips & Tricks, Odds & Ends



A solid green vertical bar is located on the left side of the slide.

# **Background:**




# **Relational Databases**



# Databases

Databases store and manipulate data.

There are many different ways to organize data...



# Relational Databases

All data is organized in Tables.

# Relational Databases

Table **Members**:

ID	Name	Place	Age
1	Joe	Sydney	34
2	Jenny	New York	42
3	Jeremy	Moskow	55

# Relational Databases

Table **Members**:

ID	Name	Place	Age
1	Joe	Sydney	34
2	Jenny	New York	42
3	Jeremy	Moskow	55

# Relational Databases

Table **Members**:

ID	Name	Place	Age
1	Joe	Sydney	34
2	Jenny	New York	42
3	Jeremy	Moskow	55

# Relational Databases

Table **Members**:

ID	Name	Place	Age
1	Joe	Sydney	34
2	Jenny	New York	42
3	Jeremy	Moskow	55

# Relational Databases

Table **Members**:

ID	Name	Place	Age
1	Joe	Sydney	34
2	Jenny	New York	42
3	Jeremy	Moskow	55

# Relational Databases

Table **Members**:

ID	Name	Place	Age
1	Joe	Sydney	34
2	Jenny	New York	42
3	Jeremy	Moskow	55



# Data Types

Fields have a type:

Text

Integer

Numeric

Date

...

# Relational Databases

Table **Members**:

ID	Name	Place	Age
1	Joe	Sydney	34
2	Jenny	New York	42
3	Jeremy	Moskow	55

# Structured Query Language

# Data Access using SQL

```
SELECT Name  
FROM Members;
```

Name
Joe
Jenny
Jeremy

# Data Access using SQL

```
SELECT Name, Age  
FROM Members;
```

Name	Age
Jenny	42
Joe	34
Jeremy	55

# Data Access using SQL

```
SELECT Name, Age  
FROM Members  
ORDER BY Age;
```

Name	Age
Joe	34
Jenny	42
Jeremy	55

# Data Access using SQL

```
SELECT Name, Age  
FROM Members  
WHERE Age > 40  
ORDER BY Name;
```

Name	Age
Jenny	42
Jeremy	55

# INSERT, UPDATE, DELETE

```
INSERT INTO Members (Name, Place, Age)  
VALUES ('Julia', 'London', 27);
```

```
UPDATE Members SET Place = 'Helsinki'  
WHERE Id = 2;
```

```
DELETE FROM Members  
WHERE Name = 'Jeremy';
```



# INSERT, UPDATE, DELETE

```
INSERT INTO Members (Name, Place, Age)  
VALUES ('Julia', 'London', 27);
```

```
UPDATE Members SET Place = 'Helsinki'  
WHERE Id = 2;
```

```
DELETE FROM Members  
WHERE Name = 'Jeremy';
```

# INSERT, UPDATE, DELETE

```
INSERT INTO Members (Name, Place, Age)  
VALUES ('Julia', 'London', 27);
```

```
UPDATE Members SET Place = 'Helsinki'  
WHERE Id = 2;
```

```
DELETE FROM Members  
WHERE Name = 'Jeremy';
```




# Advanced SQL: Aggregate Functions

```
SELECT Avg(Age)  
FROM Members;
```

Avg
43

# Advanced SQL: JOIN



# The Magic

You define a structure (“schema”).

You add data.

You ask for data back.

The database software does everything else.

# Not so magic...

Performance can depend on structure

You still need to know a bit...

# Indexes

Indexes allow faster access  
for some queries

Tradeoff: Indexes need space and need to be  
updated vs. faster queries

# Relational Databases

PostgreSQL

MySQL, MariaDB

SQLite

...



# Relational Databases

PostgreSQL

MySQL, MariaDB

SQLite

...

# PostgreSQL

Open Source

lots of features

good documentation, books, etc.

popular, great eco-system, well-supported

powerful plugin system


A solid green vertical bar is located on the left side of the slide.

# **Background:**

# **Geodata**


# Simple Feature Model

Point




MultiPoint

LineString



MultiLineString

Polygon



MultiPolygon

**We want to store this  
in a database**

# A Naïve Approach...

ID	Name	X	Y
1	Joe	151.22	-33.85
2	Jenny	-74.01	40.70
3	Jeremy	37.61	55.95

# A Better Approach...

Text

Integer

Numeric

Date

Geometry

...

# A Better Approach...

ID	Name	Geom
1	Joe	POINT(151.22 -33.85)
2	Jenny	POINT(-74.01 40.70)
3	Jeremy	POINT(37.61 55.95)



# PostGIS: Plugin

```
CREATE EXTENSION postgis;
```

# PostGIS: Datatypes

GEOMETRY

GEOMETRY (POINT)  
GEOMETRY (LINESTRING)  
GEOMETRY (POLYGON)


(also: GEOGRAPHY)



Source: Wikipedia




Source: Wikipedia

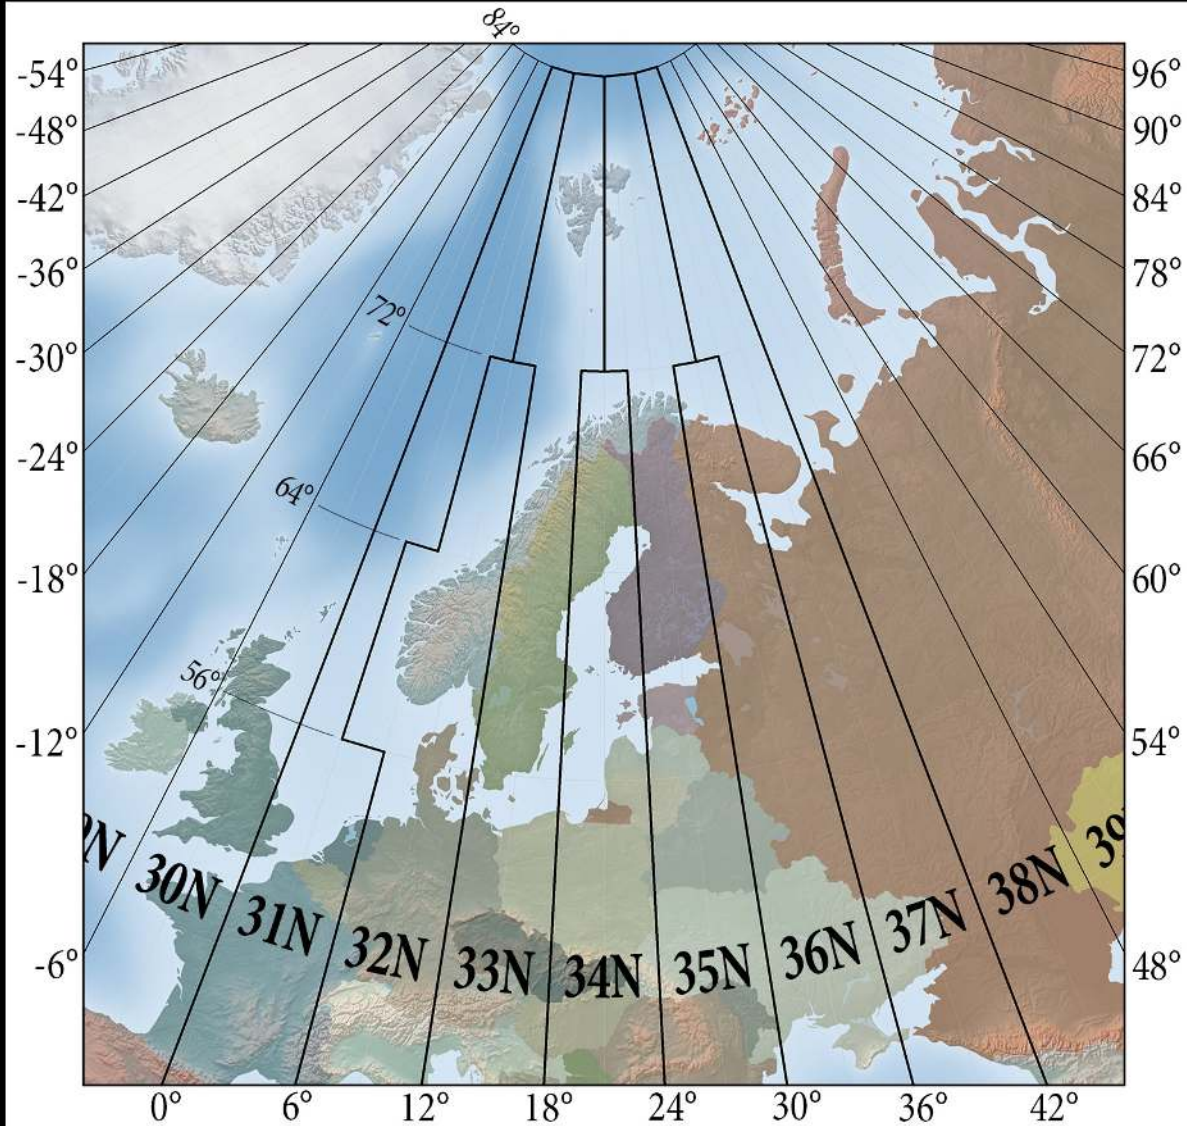


Source: Wikipedia





Source: Wikipedia



Modified UTM Zones - 31N thru 37N - Projection: Polar Stereographic

Source: Wikipedia

# PostGIS: CRS

PostGIS knows > 5000 Coordinate Systems  
(CRS/SRS)

Each Geometry associated with SRID.

Allows Transformations

Mix and match data sources



# Most important CRSes:

WGS84 – EPSG:4326

Web Mercator – EPSG:3857

# PostGIS: Datatypes

GEOMETRY (POINT, 4326)

GEOMETRY (LINESTRING, 4326)

GEOMETRY (POLYGON, 4326)

# Coordinates

Always first X axis, then Y axis  
(as in mathematics).

so: longitude first, then latitude.

# Well Known Text (WKT)

POINT(4 3)

LINESTRING(12 4, 3 2, 7, 9)

POLYGON((0 0, 4 0, 4 4, 0 4, 0 0))

MULTIPOINT / -LINESTRING / -POLYGON

# PostGIS: Indexes


Normal indexes are good for  
1-dimensional data

Spatial indexes are good for  
2/3-dimensional data  
(R -tree)

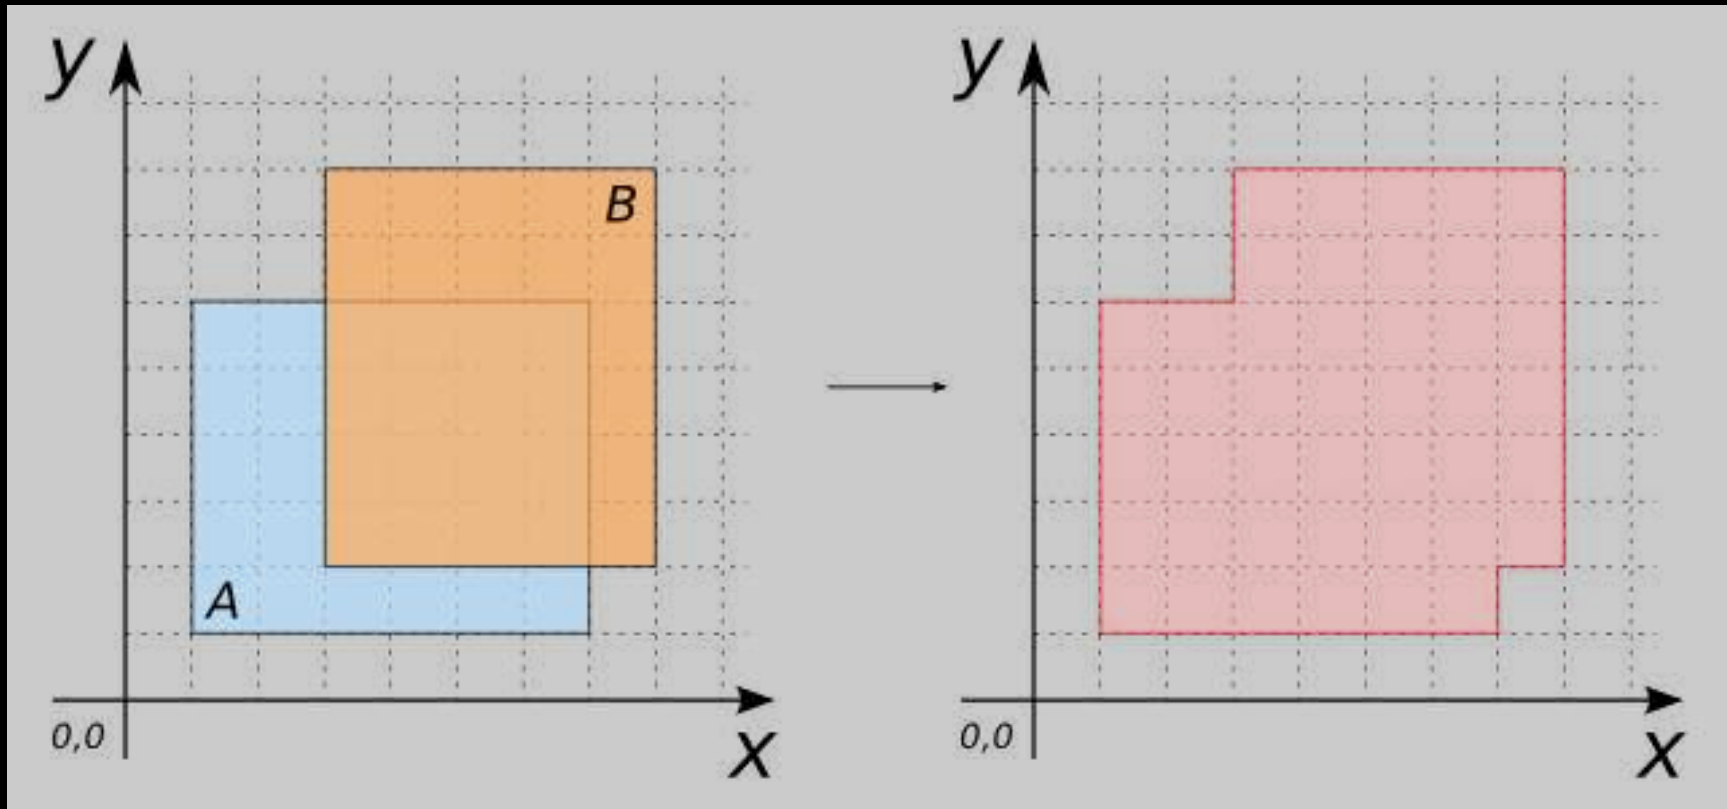
# PostGIS: Operations

Huge number of  
operations on spatial data

# PostGIS: ST\_Contains




# PostGIS: ST\_Union







# PostGIS: ST\_Intersection




# PostGIS: ST\_Buffer



`ST_Buffer(geom, 1)`  
(by default `join=round`)




`ST_Buffer(geom, 1, 'join=mitre')`  
`ST_Buffer(geom, 1, 'join=miter')`



`ST_Buffer(geom, 1, 'join=bevel')`

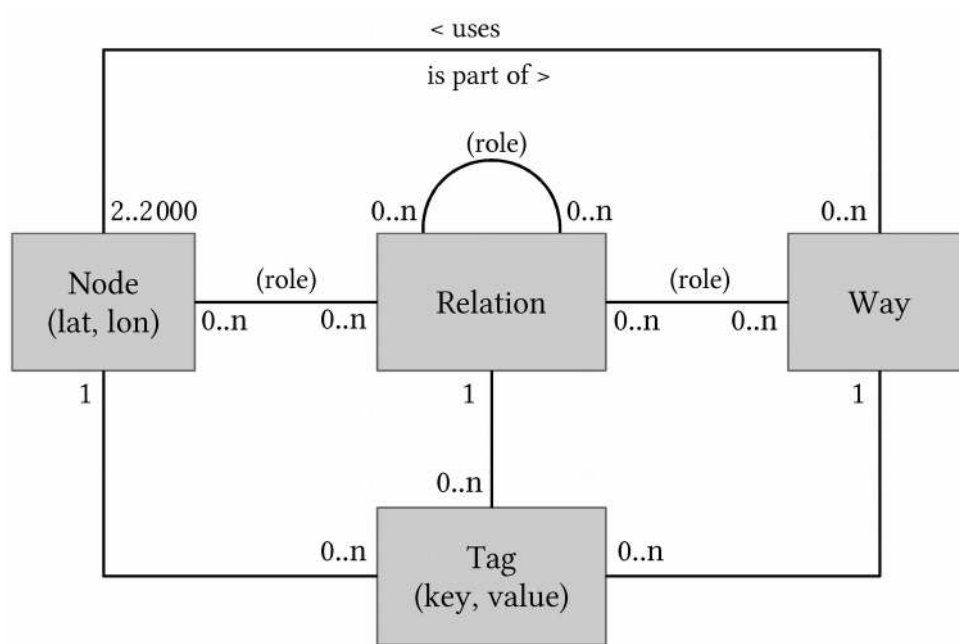
# PostGIS: ST\_ShortestLine



A solid green vertical bar is located on the far left side of the image, extending from the top to the bottom.

# Converting OSM Data

# OSM Data Model



# Mismatch

OSM  
Data Model

Relational /  
Simple Feature  
Data Model


# Mismatch

OSM  
Data Model





Conversion

Relational /  
Simple Feature  
Data Model







# Conversion: Selection

What data do we actually need?

nodes, ways, relations?

user id, timestamp, version, ...?

which tags?

# Conversion: Data Types

tags in OSM:  
key → value (both text)

Map to:  
text, integer, boolean, enums, ...

# Conversion: Tags → Attributes

tags in OSM are flexible, table columns are fixed

Highway	Name	Oneway
primary	Main St	false
residential	Elm St	true
trunk		true

# Conversion: Tags → Attributes

tags in OSM are flexible, table columns are fixed

Highway	Name	Oneway
primary	Main St	false
residential	Elm St	true
trunk		true

# Conversion: Tags → Attributes

tags in OSM are flexible, table columns are fixed

Highway	Name	Oneway
primary	Main St	false
residential	Elm St	true
trunk		true

# Conversion: hstore and JSON

Place	Name
city	de: München, en: Munich
city	de: Aachen, fr: Aix-la-Chapelle
village	de: Lübben, hsb: Lubin

# Conversion: Tables

split data into tables...

few tables vs. many tables

by geometry type and/or by subject type



# Split by Geometry Type

Tables:

nodes

ways

areas

# Split by Geometry/Feature Class

Tables:

restaurants  
bus\_stops  
addresses  
places

...

highways  
railways  
rivers  
powerlines

...

lakes  
forests  
countries  
buildings

...

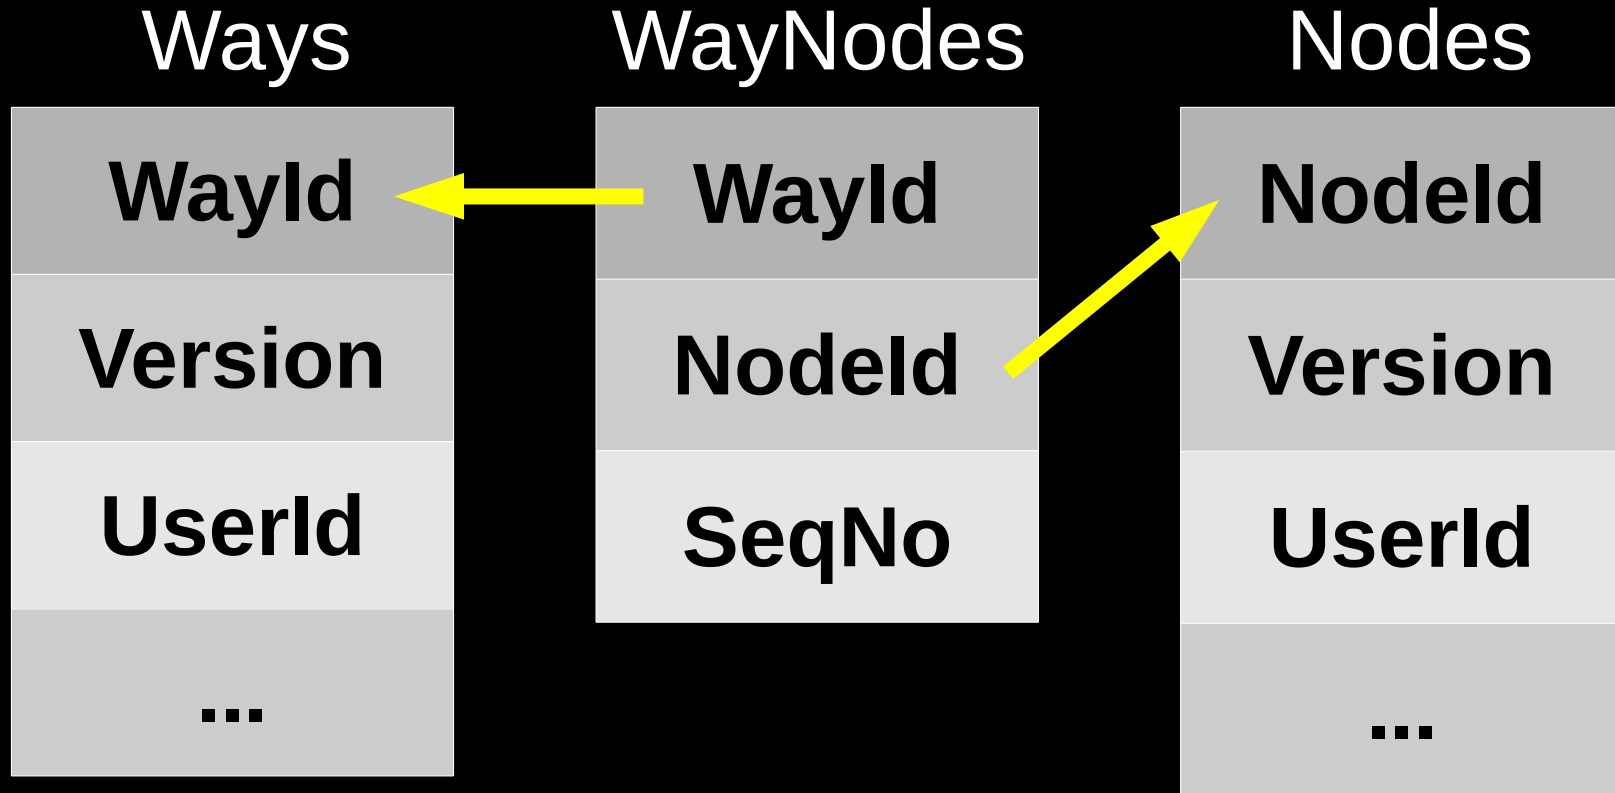
# Conversion: Handling lists

nodes in ways


members in relations

tags in nodes, ways, or relations

# Conversion: Way Nodes



# Conversion: Way Nodes



# Conversion: Relation Members

similar to way nodes

but

array of tuple (type, id, role)

# Conversion: Geometry

Nodes → Points

Ways → LineStrings / Polygons

Multipolygon relations → Polygons

Route relations → MultiLineStrings

...

# Conversion: Geometry

Generalized geometries

For lower zoom levels / small scales



Selection – Merging – Simplification






# Conversion: Where?

Conversion can happen

1. before import in code
2. after import in the DB





# Conversion: Where?

Conversion can happen

1. before import in code
2. after import in the DB

# Conversion: Where?

Conversion can happen

**fast**

1. before import in code

2. after import in the DB

**flexible**

# Conversion: Assemble Lines


Take Locations from Nodes

Assemble them into LineStrings

# Conversion: Node locations

Where to store node locations?

1. in the database
2. in specialized index




Node Location Store



# Conversion: Polygons

Assemble (Multi)Polygons


from Ways/Relations




lots of trade-offs

not all software/schemas support updates

# Complete data needed for updates



# Complete data needed for updates



# Complete data needed for updates

Two kinds of data:

1. The data you need for you application
2. The data needed to allow updating

# Complete data needed for updates

Two kinds of data:

1. The data you need for you application
2. The data needed to allow updating

Where? Database? External Storage?

# Snapshot vs. History

most use cases only need  
**current** OSM data

some need  
**history** of OSM data

# Snapshot vs. History

most use cases only need  
**current** OSM data


some need  
**history** of OSM data



**Much more effort needed !**



A solid green vertical bar is located on the far left side of the image.

# Use Cases





# Use Cases

API DB

Rendering

Geocoding

Routing

Analytics

# API DB

Schema used in the main OSM database

PostgreSQL - No PostGIS !

Normal access via HTTP API

You can run your own

# API DB

Needs all (also historical) data

Multiple writers, transactions

Allow bounding-box download

Allow read/write access

Create full dumps and replication diffs

# Rendering

Turning data into maps

Render into bitmap, vector tiles, etc.

# Rendering

Get all data for an area quickly

Multiple layers

Create generalized geometries

One writer, multiple reader

Regular updates



# Geocoding

“Search”

Geocoding – Address to Location

Reverse Geocoding – Location to Address

# Geocoding

Build address hierarchy

Quick “fuzzy” search

One writer, multiple reader

Regular updates

# Routing

Using PostgreSQL plugin PgRouting

Flexible, but slow

# Routing

Build network of streets

Calculate weights

Find route through network

# **Analytics**

Statistics

Comparing data

Conflating data

Many diverse needs

# Example 1: Wind Power

Find a place that ...

- ... has lots of steady winds
- ... is near existing high voltage lines
- ... is far from residential areas

# Example 2: Public Transport

How far is the nearest public transport stop?

How many people live where the nearest stop is more than  $x$  meters away?

Where should a new bus route go?

# Example 3: OSM Contributors

Who are the most active OSM contributors?

What kinds of things do they map?

Where do they do their mapping?



# **Analytics**

Flexible data model

Use of many geometric operations

Batch processing in multiple steps

A solid green vertical bar is located on the far left side of the image.

# Tools

# psql

```
File Edit View Terminal Tabs Help
sqrt=# select osm_id, osm_type, highway, substr(ST_AsText(geom), 1, 46) as geom
from highway_motorway limit 10;
 osm_id | osm_type | highway | geom
-----+-----+-----+-----
3700602 | way      | motorway | LINESTRING(12.3957525 51.4045029,12.3940079 51.
4068262 | way      | motorway | LINESTRING(12.1950993 51.4221162,12.195996 51.
4068272 | way      | motorway | LINESTRING(12.1956903 51.4236168,12.194645 51.
4068276 | way      | motorway | LINESTRING(12.2079569 51.426431,12.2111331 51.
4071131 | way      | motorway | LINESTRING(12.375885 51.4084059,12.3754508 51.
4071132 | way      | motorway | LINESTRING(12.355407 51.4120578,12.353102 51.4
4071145 | way      | motorway | LINESTRING(12.2989367 51.4155116,12.2967172 51
4071159 | way      | motorway | LINESTRING(12.1805927 51.3958729,12.1805298 51
4100106 | way      | motorway | LINESTRING(12.1999562 51.4310166,12.1979877 51
4259290 | way      | motorway | LINESTRING(12.5483554 51.3050722,12.5492398 51
(10 rows)

sqrt=#
```

# psql

```
File Edit View Terminal Tabs Help
sqrt=# select osm_id, osm_type, highway, substr(ST_AsText(geom), 1, 46) as geom
from highway_motorway limit 10;
 osm_id | osm_type | highway |          geom
-----+-----+-----+-----
 3700602 | way      | motorway | LINESTRING(12.3957525 51.4045029,12.3940079 51.
 4068262 | way      | motorway | LINESTRING(12.1950993 51.4221162,12.195996 51.
 4068272 | way      | motorway | LINESTRING(12.1956903 51.4236168,12.194645 51.
 4068276 | way      | motorway | LINESTRING(12.2079569 51.426431,12.2111331 51.
 4071131 | way      | motorway | LINESTRING(12.375885 51.4084059,12.3754508 51.
 4071132 | way      | motorway | LINESTRING(12.355407 51.4120578,12.353102 51.4
 4071145 | way      | motorway | LINESTRING(12.2989367 51.4155116,12.2967172 51
 4071159 | way      | motorway | LINESTRING(12.1805927 51.3958729,12.1805298 51
 4100106 | way      | motorway | LINESTRING(12.1999562 51.4310166,12.1979877 51
 4259290 | way      | motorway | LINESTRING(12.5483554 51.3050722,12.5492398 51
(10 rows)

sqrt=#
```



# Editor

# pgadmin

pgAdmin 4

File Object Tools Help

Browser

- geo\_test
  - Casts
  - Catalogs
  - Event Triggers
  - Extensions
  - Foreign Data Wrappers
  - Languages
  - Schemas (2)
    - public
      - Collations
      - Domains
      - FTS Configurations
      - FTS Dictionaries
      - FTS Parsers
      - FTS Templates
      - Foreign Tables
      - Functions
      - Materialized Views
      - Sequences
      - Tables (6)
        - boundary\_claims
        - city\_points
        - country\_outlines
        - random\_geometries
        - random\_geometries
        - spatial\_ref\_sys
      - Trigger Functions
      - Types
      - Views
      - topology

Dashboard Properties SQL Statistics Dependencies Dependents Edit Data - pem ... Edit Data - geo\_test on postgres@PostgreSQL

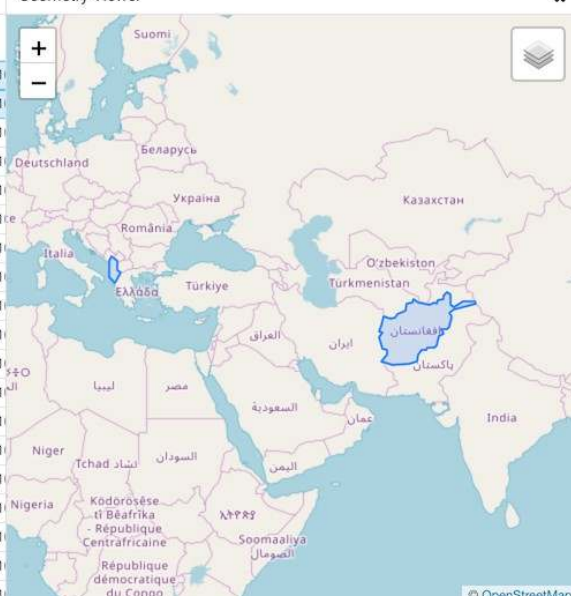
Query Editor Query History

```
1 SELECT * FROM public.country_outlines
2
```

Data Output Explain Messages Notifications




	ogc_fid [PK] integer	id character varying	name character varying	geometry4326 geometry
1	1	AFG	Afghanistan	0103000020E61
2	3	ALB	Albania	0103000020E61
3	7	ATA	Antarctica	0106000020E61
4	18	BHS	The Bahamas	0106000020E61
5	6	ARM	Armenia	0103000020E61
6	314	QAT	Qatar	0103000020E61
7	365	ATA	Antarctica	0106000020E61
8	269	KOR	South Korea	0103000020E61
9	26	BTN	Bhutan	0103000020E61
10	65	GNQ	Equatorial Guinea	0103000020E61
11	48	ECU	Ecuador	0103000020E61
12	8	ATF	French Southern and...	0103000020E61
13	449	CS-KM	Kosovo	0103000020E61
14	92	KWT	Kuwait	0103000020E61
15	19	BIH	Bosnia and Herzego...	0103000020E61
16	225	DOM	Dominican Republic	0103000020E61
17	131	PRI	Puerto Rico	0103000020E61
18	20	BLR	Belarus	0103000020E61
19	325	SLV	El Salvador	0103000020E61

Geometry Viewer



© OpenStreetMap





# Osmosis

<https://wiki.osm.org/wiki/Osmosis>

Use case: API DB, Analytics

Updates: Yes

Schema: Several

Status: Not being maintained



# Osmosis Schemas

API DB (version 0.6)

PostGIS Snapshot Schema (uses hstore)

PostGIS Simple Schema (no hstore)

(API DB MySQL <0.6)

# osm2pgsql

Used in “standard” OSM rendering toolchain

<https://wiki.osm.org/wiki/Osm2pgsql>

Use case: Rendering

Updates: Yes

Schema: Few tables (hstore optional)

Status: Maintained

# Imposm3

Alternative rendering toolchain

<https://imposm.org>

Use case: Rendering

Updates: Yes

Schema: Many tables

Status: Actively maintained

# Nominatim

Standard OSM search/geocoding

<https://nominatim.org>

Use case: (Reverse) Geocoding

Schema: Optimized for geocoding

Status: Actively maintained

Uses osm2pgsql (with special plugin)

# Osmium

<https://osmcode.org/osmium-tool/>

Use case: Analytics, (Rendering)

Updates: No

Schema: Simple

Status: Actively maintained

Simple to run for ad-hoc use

# osm-postgresql-experiments

Experimental, very flexible data import

<https://github.com/osmcode/osm-postgresql-experiments>

Use case: Rendering, Analytics

Updates: (Yes)

Schema: Flexible

Status: Experimental

# osm2pgrouting

Importer für PgRouting




<https://github.com/pgrouting/osm2pgrouting>

Use case: Routing

Updates: No

Schema: PgRouting

Status: Maintained

A solid green vertical bar is positioned on the far left side of the image, extending from the top to the bottom.

# **Tips & Tricks**

# **Odds & Ends**



# Quantity Structure

How much disk space do I need?

How much memory do I need?

How long will an import take?

# Quantity Structure

hundreds of Gbytes  
for full planet

How much disk space do I need?

How much memory do I need?

How long will an import take?

# Quantity Structure

hundreds of Gbytes  
for full planet

How much disk space do I need?

How much memory do I need?

More!

How long will an import take?

# Quantity Structure

hundreds of Gbytes  
for full planet

How much disk space do I need?

How much memory do I need?

More!

How long will an import take?

many hours  
if not days for planet

# Start small...

Do not try to import the whole planet at first!

Start small  
(e.g. with data for a city)  
and  
work your way up

# Minimize data

1. Filter data outside DB if you can
2. Import data into DB

# Importing Data

1. load data

2. create indexes

3. ANALYZE

# Performance Tuning

You **will** need to tune your PostgreSQL!

Settings in postgresql.conf:

shared\_buffers, work\_mem, maintenance\_work\_mem, fsync,  
synchronous\_commit, checkpoint\_timeout,  
checkpoint\_completion\_target, ...



# Indexes

Learn how indexes work and  
when they are used

Also for spatial indexes!

Use EXPLAIN command

# The COPY command

COPY instead of INSERT

more efficient, use it if possible

# Learning Curve

PostgreSQL / PostGIS  
is an incredible powerful tool

“Magic” working of PostgreSQL  
can be surprising

Configure the logs and look at them

# Learning Curve

Spatial operations are extra magic!

Again: Start small

# Always growing

Database will grow over time  
(not only because of more OSM data)

VACUUM

Still grow more

# Other SQL Databases

MySql, MariaDB

Oracle Spatial

Sqlite (Library, not Server)

# Other SQL Databases

geodata support lacking

MySQL, MariaDB

Oracle Spatial

Sqlite (Library, not Server)

# Other SQL Databases

geodata support lacking

MySQL, MariaDB

Oracle Spatial

proprietary

Sqlite (Library, not Server)



# Other SQL Databases

geodata support lacking

MySql, MariaDB

Oracle Spatial

proprietary

Sqlite (Library, not Server)

not as powerful, problems with huge datasets,  
but can be useful for some applications



# THE END

Jochen Topf  
jochentopf.com  
jochen@topf.org